



**POLITECNICO
DI TORINO**

Recap Big data: architectures and data analytics (01QYDOV)

Jacopo Nasi
Computer Engineer
Politecnico di Torino

II Period - 2018/2019

April 1, 2019

Contents

1	Introduction	4
2	Apache Hadoop	4
2.1	Structure of a program	5
2.2	Data Types	6
2.3	Combiner and Counters	7
2.4	Configurations	7
2.5	Map-only job	8
2.6	Setup and Cleanup methods	8
2.7	Patterns	8
2.8	Multiple Inputs/Outputs	9
2.9	Distributed Cache	10
2.10	Data Organization Patterns	10
2.11	Metapatterns	10
2.12	Join Patterns	11

License

This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 3.0 Unported License.

You are free:

- **to Share:** to copy, distribute and transmit the work
- **to Remix:** to adapt the work

Under the following conditions:

- **Attribution:** you must attribute the work in the manner specified by the author or licensor (but not in any way that suggests that they endorse you or your use of the work)
- **Noncommercial:** you may not use this work for commercial purposes.
- **Share Alike:** if you alter, transform, or build upon this work, you may distribute the resulting work only under the same or similar license to this one.

More information on the Creative Commons website (<http://creativecommons.org>).



Acknowledgments

Questo breve riepilogo non ha alcuno scopo se non quello di agevolare lo studio di me stesso, se vi fosse di aiuto siete liberi di usarlo.

Le fonti su cui mi sono basato sono quelle relative al corso offerto (**Big data: architectures and data analytics (01QYDOV)**) dal Politecnico di Torino durante l'anno accademico 2018/2019.

Non mi assumo nessuna responsabilità in merito ad errori o qualsiasi altra cosa. Fatene buon uso!

1 Introduction

The amount of data increases every days, to analyze them, system must scale according. All these systems must deal with computation capabilities, network bandwidth and failures. A good practice to solve problems is the parallelizations, a single-node architecture cannot solve, in reasonable time, the request of these days. For these reasons, nowadays, clusters (fig.1) are pretty common.




Figure 1: Cluster Architecture

The main characteristics are:

- **Scalability:** Increase number of users, amount of datas, complexity of problems
 - Scale **UP**: More power/resources for a single-node
 - Scale **OUT**: Add more nodws to the cluster

2 Apache Hadoop

Scalable fault-tolerant distributed system for Big Data: Distributed Data Storage, Distributed Data Processing, etc... Is an Open Source project under Apache license. Is used by a lot of main IT companies all around the world. The core of the structure are:

- **Distributed BD Processing Infrastructure** (MapReduce):

- High-level abstraction view: Not require to take care of scheduling and synchronizations
- Fault Tolerant: Node and task are automatically managed
- **HDFS** (Hadoop Distributed File System):
 - Fault-Tolerant
 - High availability distributed storage

The paradigm of MapReduce will not be illustrated, in figure 2 a sample of the process during the word count problem.

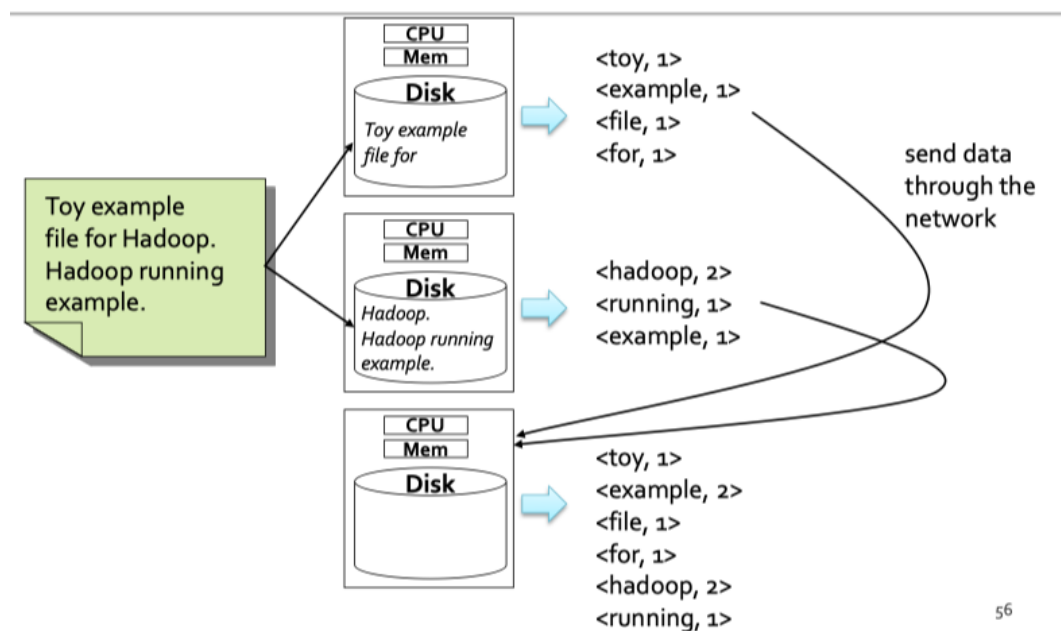


Figure 2: MapReduce Example

2.1 Structure of a program

The project is divided in different parts:

- Driver: Coordinators of the jobs
- Mapper: Class for the Map part
- Reducer: Class implementing the Reduce part

There are also other important terms:

- Job: Execution of a MapReduce code over a dataset
- Task: Execution of a Map task or a Reducer task on a slice of data

The workflow of a project is similar to the following one:




Figure 3: MapReduce Workflow in Hadoop

2.2 Data Types

Hadoop make use of its own basic data types:

- Text: Like Java string
- IntWritable: Like java integer
- LongWritable
- FloatWritable

Also the data for the input and output has its own type:

- TextInputFormat (Lines)
- KeyValueTextInputFormat (2 Fields)
- SequenceFileInputFormat (Binary)

Also the output have similar formats.

Is also possible to define personalized data types, these also the usage of more complex values. To use them, is required, to implements the *org.apache.hadoop.io.Writable* interface and its methods, in particular:

- write()
- readFields()

also the method toString is important, in order to generate the value to be emit with the pair.

2.3 Combiner and Counters

In standard MapReduce applications the (key,value) pairs emitted by mappers are sent to the reducers through the network. In some case, a sort of "pre-aggregations" could be useful to improve performances and reduce the amount of data shared on the network.

An example with the word count problem with clarify the usage of combiners.

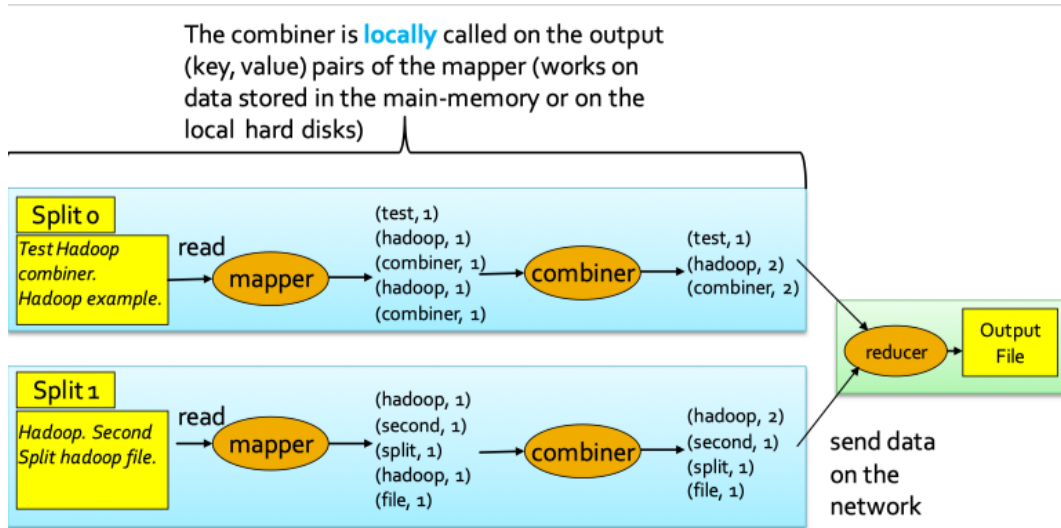


Figure 4: Example of combiner usage

Is an instance of the Reducer class, that implements a pre-reduce phase, characterized by the `reduce()` method. It process (key, [list of value]) and emits (key, value). It runs on the cluster.

Combiners can be also exploited by in-mapper solutions, by initialize a set of in-mapper variables during the instance of the Mapper. They are created by using `setup/cleanup` methods.

Hadoop provides a set of basic, built-in, counters to store some statistics about jobs, mappers and reducers. Also some ad-hoc, user-defined, counters can be defined. They can be used by the exploiting enum Java types, and increment by using the following code:

```
context.getCounter(countername).increment(value);
```

2.4 Configurations

Is possible to define some properties shared among the different part (M,C,R), they must be defined in the driver using the following syntax:

```
conf.set("property-name", "value");
```

and they can be retrieved in the Mapper and/or Reducer phase by using:

```
context.getConfiguration().get("property-name");
```

2.5 Map-only job

In some cases everything can be achieved using the Mapper phase only. Of course, Hadoop, allows the execution of this kind of jobs. Everything is achieved by setting the number of reducers task to zero. **Not setting this value can generate problems during the executions.**

2.6 Setup and Cleanup methods

Are two optional methods to be invoked in the Mapper class, one at the start (Setup()) and the other at the end, both invoked one time for each mapper.

2.7 Patterns

Common patterns to solve problems. The following are known as **Summarization Patterns**:

Numerical Summarizations group records/objects by a key field and calculate a numerical aggregate (avg, min, max, std deviations, etc...). Are used to provide a top-level view of large input data sets. They are used to analyze, by domain experts, trends or anomalies.

The phase:

- **Mappers:** Key is used to define groups, Value for computing aggregate statistics
- **Reducers:** Receive a set of numerical values for each "group-by" key and compute the final stats
- **Combiners:** Used to speed-up the process

Inverted Index Summarizations the goal is to build an index from the input data to support faster search or data enrichment.

- **Mappers:** Key is the set of fields to index (a keyword), Value is a unique identifier of the objects to associate with each "keyword"
- **Reducers:** Receive a set of identifiers for each keyword and simply concatenate them
- **Combiners:** Normally NOT used

Counting with Counters the goal is to compute count summarizations of data sets. Also this solutions can be used to identify trends and anomalies.

- **Mappers:** Process each input record and increment a set of counters

This pattern can be achieved by using also MapOnly jobs, the result are directly printed by the Driver of the application.

The following are know as **Filtering Patterns**:

Filtering has the goal to filter out input records that are not of interest and keep the others, in order to focusing only on the records of interest. The pourpose is to reduce the set of the data for the analysis.

- **Mappers:** Output one pair for each record that satisfies the enforced rule
- **Reducers:** Almost useless

Top K is select a small set of top K record according to a ranking function. Used to focusing only on the most important (for the ranking function) values.

- **Mappers:** Each mapper initializes an in-mapper top K list (small), cleanup emits the pairs after the complete scan
- **Reducers:** a Single Instance compute the final top k list by merging the local lists

Distinct it finds a unique set of values/records, prune the duplicates.

- **Mappers:** Emit one (key, null) pair for each input record
- **Reducers:** Emit one (key, null) pair for each input (key, [list of values]) pair

2.8 Multiple Inputs/Outputs

Hadoop allows the user to implements different input and/or outputs, it could be useful in some case. The input differentiation is achieved by using the addInputPath() method, the output one is achieved with addNamedOutput() method.

2.9 Distributed Cache

Some application need to share and cache (small) read-only files to perform efficiently their tasks. These files should be accessible by all nodes of the cluster in an efficient way. The Distributed Cache is a facility provided by the Hadoop-based MapReduce framework to cache files. These files are specified in the Driver, during the initialization of the job, Hadoop create a local copy of the shared/cached files in all nodes which are used to execute some tasks of the job, the cache file is then read by the Mapper, usually in the setup method. Save with:

```
job.addCacheFile(new Path("hdfsPath").toUri());
```

And it is retrieved by using:

```
URI[] urisCachedFiles = context.getCacheFiles();
```

2.10 Data Organization Patterns

Shuffling is a technique used to randomize the order of the data, the motivations are:

- Anonymization of data
- Selecting a subset of random data

The structure is the following:

- **Mappers:** Emit one (key (random), value) pair for each input record
- **Reducers:** Emit one (key, null) pair for each value in [list of values] of the input (key, [list of values]) pair

2.11 Metapatterns

Job Chaining is the technique to execute a sequence of jobs (synchronizing them), the intent is to manage workflow of complex applications based on many phases.

The driver (single instance) contains the whole workflow and it execute the jobs in the proper order.




Figure 5: Job Chaining

2.12 Join Patterns

Reduce side natural join The goal is to join the content of two relations (both large), this operation is useful in many applications.

Suppose you want to join the following tables:

- **Users:** userid, name, surname
- **Likes:** userid, movieGenre

The record of the first table will be:

(userid=u1, "Users:name=Paolo,surname=Garza")

While the ones of the second table will be:

(userid=u1, "Likes:movieGenre=horror") (userid=u1,
"Likes:movieGenre=adventure")

The reducers iterate over the values associated with each key (value of the common attributes) and compute the "local natural join" for the current key. The output will be like the followings:

- (userid=u1, "Users:name=Paolo,surname=Garza,movieGenre=horror")
- (userid=u1, "Users:name=Paolo,surname=Garza,movieGenre=adventure")

Map side natural join the goal is to join the content of the two relations, with a large and a small table (can be loaded in main memory). Is a **Map-only** job where the mapper process the content of the large table, the distributed cache is used to provide a copy of the small table to all mappers. The process is than a "local natural join" like in the previous pattern. Also the other kind of joins can be achieved.